

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn

In [2]: import glob
path="C:\\S1_Dataset"
all_files=glob.glob(path+'/*')
print(all_files)

['C:\\S1_Dataset\\d1p01M', 'C:\\S1_Dataset\\d1p02M', 'C:\\S1_Dataset\\d1p03M', 'C:\\S1_Dataset\\d1p04M', 'C:\\S1_Dataset\\d1p05M', 'C:\\S1_Dataset\\d1p06M',
'C:\\S1_Dataset\\d1p07M', 'C:\\S1_Dataset\\d1p08F', 'C:\\S1_Dataset\\d1p09F', 'C:\\S1_Dataset\\d1p10F', 'C:\\S1_Dataset\\d1p11F', 'C:\\S1_Dataset\\d1p12F',
'C:\\S1_Dataset\\d1p13F', 'C:\\S1_Dataset\\d1p14F', 'C:\\S1_Dataset\\d1p15F', 'C:\\S1_Dataset\\d1p16F', 'C:\\S1_Dataset\\d1p17F', 'C:\\S1_Dataset\\d1p18F',
'C:\\S1_Dataset\\d1p19F', 'C:\\S1_Dataset\\d1p20F', 'C:\\S1_Dataset\\d1p21F', 'C:\\S1_Dataset\\d1p22F', 'C:\\S1_Dataset\\d1p23F', 'C:\\S1_Dataset\\d1p24F',
'C:\\S1_Dataset\\d1p25F', 'C:\\S1_Dataset\\d1p26F', 'C:\\S1_Dataset\\d1p27F', 'C:\\S1_Dataset\\d1p28F', 'C:\\S1_Dataset\\d1p29F', 'C:\\S1_Dataset\\d1p30F',
'C:\\S1_Dataset\\d1p31F', 'C:\\S1_Dataset\\d1p32F', 'C:\\S1_Dataset\\d1p33F', 'C:\\S1_Dataset\\d1p34F', 'C:\\S1_Dataset\\d1p35F', 'C:\\S1_Dataset\\d1p36M',
'C:\\S1_Dataset\\d1p37M', 'C:\\S1_Dataset\\d1p38M', 'C:\\S1_Dataset\\d1p39M', 'C:\\S1_Dataset\\d1p40M', 'C:\\S1_Dataset\\d1p41M', 'C:\\S1_Dataset\\d1p42M',
'C:\\S1_Dataset\\d1p43M', 'C:\\S1_Dataset\\d1p44M', 'C:\\S1_Dataset\\d1p45M', 'C:\\S1_Dataset\\d1p46M', 'C:\\S1_Dataset\\d1p47M', 'C:\\S1_Dataset\\d1p48M',
'C:\\S1_Dataset\\d1p49F', 'C:\\S1_Dataset\\d1p50F', 'C:\\S1_Dataset\\d1p51F', 'C:\\S1_Dataset\\d1p52F', 'C:\\S1_Dataset\\d1p53F', 'C:\\S1_Dataset\\d1p54F',
'C:\\S1_Dataset\\d1p55F', 'C:\\S1_Dataset\\d1p56F', 'C:\\S1_Dataset\\d1p57F', 'C:\\S1_Dataset\\d1p58F', 'C:\\S1_Dataset\\d1p59F', 'C:\\S1_Dataset\\d1p60F']

In [3]: li=[]
for file in all_files:
    #print(file)
    if file.endswith('.txt'):
        continue
    df=pd.read_csv(file,header=None,index_col=None)
    li.append(df)
healthy_op=pd.concat(li,axis=0,ignore_index=True)

In [4]: healthy_op.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 52482 entries, 0 to 52481
Data columns (total 9 columns):
#   Column   Non-Null Count  Dtype
---  ------  -
0    0         52482 non-null  float64
1    1         52482 non-null  float64
2    2         52482 non-null  float64
3    3         52482 non-null  float64
4    4         52482 non-null  int64
5    5         52482 non-null  float64
6    6         52482 non-null  float64
7    7         52482 non-null  float64
8    8         52482 non-null  int64
dtypes: float64(7), int64(2)
memory usage: 3.6 MB

In [5]: healthy_op.columns=['Time','Acc.Front','Acc.vert','Acc.Lat','id','RSSI','Phase','Freq','Activity_Label']
healthy_op.head()

Out[5]:
   Time  Acc.Front  Acc.vert  Acc.Lat  id  RSSI  Phase  Freq  Activity_Label
0  0.00   0.27203   1.00820 -0.082102   1  -63.5  2.4252  924.25           1
1  0.50   0.27203   1.00820 -0.082102   1  -63.0  4.7369  921.75           1
2  1.50   0.44791   0.91636 -0.013684   1  -63.5  3.0311  923.75           1
3  1.75   0.44791   0.91636 -0.013684   1  -63.0  2.0371  921.25           1
4  2.50   0.34238   0.96229 -0.059296   1  -63.5  5.8920  920.25           1

In [6]: healthy_op.isnull().values.any()

Out[6]: False

In [7]: cols=len(healthy_op.columns)-1
healthy_op1=healthy_op.values
X=healthy_op1[:, :8]
Y=healthy_op1[:,8]

In [8]: from sklearn.preprocessing import Normalizer
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split,GridSearchCV
from sklearn.metrics import accuracy_score
from sklearn.neighbors import KNeighborsClassifier

In [9]: normalize=Normalizer()
X=normalize.fit_transform(X)
print(X)

[[ 0.00000000e+00  2.93631563e-04  1.08825991e-03 ... -6.85424559e-02
  2.61778211e-03  9.97643541e-01]
 [ 5.41175984e-04  2.94432206e-04  1.09122725e-03 ... -6.81881739e-02
  5.12699303e-03  9.97657926e-01]
 [ 1.61998019e-03  4.83736884e-04  9.89656697e-04 ... -6.85791613e-02
  3.27354797e-03  9.97637799e-01]
 ...
 [ 4.64054952e-01  3.16054398e-04  9.08811718e-04 ... -5.78276511e-02
  4.78430619e-03  8.83902815e-01]
 [ 4.65829320e-01  1.14107909e-04  7.86629014e-04 ... -5.34284525e-02
  4.68329008e-03  8.83239105e-01]
 [ 4.67376010e-01  6.95187748e-05  1.26064817e-03 ... -5.40291449e-02
  9.91649969e-04  8.82396344e-01]]

In [10]: healthy_op.shape

Out[10]: (52482, 9)

In [11]: train, test = train_test_split(healthy_op, test_size = 0.25)
print(train.shape)
print(test.shape)

(39361, 9)
(13121, 9)

In [12]: X_train = train[['Time','Acc.Front','Acc.vert','Acc.Lat','id','RSSI','Phase','Freq']]
Y_train =train.Activity_Label
X_test = test[['Time','Acc.Front','Acc.vert','Acc.Lat','id','RSSI','Phase','Freq']]
Y_test =test.Activity_Label

In [13]: X_train.head()

Out[13]:
   Time  Acc.Front  Acc.vert  Acc.Lat  id  RSSI  Phase  Freq
27023  13.850    0.18995  1.008200  0.054735  4  -57.0  0.62433  925.25
42145  434.750    1.16320  0.009187 -0.070699  1  -56.0  5.12500  922.25
48453  309.750    1.15140  0.020670 -0.104910  1  -56.5  5.69410  921.25
21562   76.175    1.04590  0.020670 -0.161920  3  -61.5  5.64810  923.25
24070  335.100    1.05760  0.020670 -0.150520  1  -55.5  5.12500  920.75

In [14]: Y_train.head()

Out[14]:
27023    1
42145    3
48453    3
21562    3
24070    3
Name: Activity_Label, dtype: int64

In [15]: #normalization
normalize=Normalizer()
n_X_train=normalize.fit_transform(X_train)

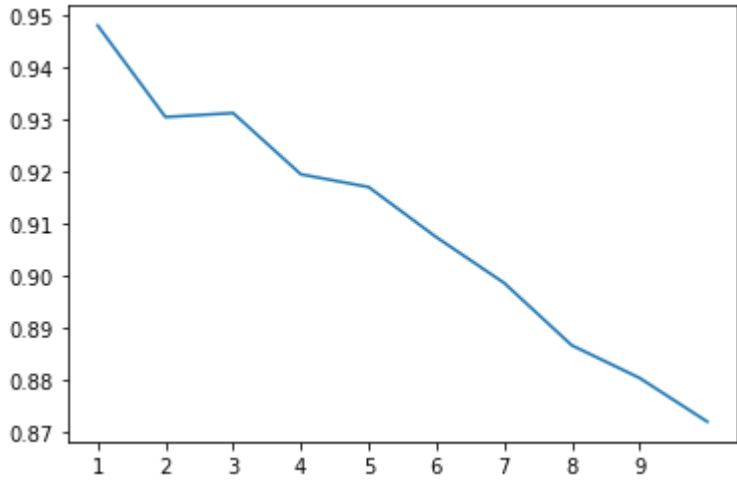
normalize=Normalizer()
n_X_test=normalize.fit_transform(X_test)

In [16]: # KNN
knn=KNeighborsClassifier(n_neighbors=3)
knn.fit(n_X_train,Y_train)
y_pred_knn=knn.predict(n_X_test)
print("accuracy KNN= ",accuracy_score(Y_test,y_pred_knn))

accuracy KNN=  0.9312552396920967

In [17]: a_index=list(range(1,11))
a=pd.Series()
x=[1,2,3,4,5,6,7,8,9]
for i in list(range(1,11)):
    model=KNeighborsClassifier(n_neighbors=1)
    model.fit(n_X_train,Y_train)
    prediction=model.predict(n_X_test)
    a=a.append(pd.Series(accuracy_score(prediction,Y_test)))
plt.plot(a_index, a)
plt.xticks(x)

C:\Users\Ganesh\AppData\Local\Temp\ipykernel_10616\3442830369.py:2: DeprecationWarning: The default dtype for empty Series will be 'object' instead of 'float64' in a future version. Specify a dtype explicitly to silence this warning.
a=pd.Series()

Out[17]: ([<matplotlib.axis.XTick at 0x27e9141da60>,
<matplotlib.axis.XTick at 0x27e9141da30>,
<matplotlib.axis.XTick at 0x27e9140d9d0>,
<matplotlib.axis.XTick at 0x27e92dec640>,
<matplotlib.axis.XTick at 0x27e92dec90>,
<matplotlib.axis.XTick at 0x27e92decdc0>,
<matplotlib.axis.XTick at 0x27e92df8640>,
<matplotlib.axis.XTick at 0x27e92df8d90>,
<matplotlib.axis.XTick at 0x27e92dfd520>],
[Text(0, 0, ''),
Text(0, 0, ''),
Text(0, 0, ''),
Text(0, 0, ''),
Text(0, 0, ''),
Text(0, 0, ''),
Text(0, 0, ''),
Text(0, 0, ''),
Text(0, 0, ''),
Text(0, 0, '')])


In [18]: lr=LogisticRegression()
estimator={'solver':['newton-cg','liblinear','lbfgs','sag']}
gsc=GridSearchCV(lr,estimator)
gsc.fit(n_X_train,Y_train)
y_gsc_pred=gsc.predict(n_X_test)
print("accuracy gsc= ",accuracy_score(Y_test,y_gsc_pred))
print(gsc.best_estimator_)

C:\Users\Ganesh\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
accuracy gsc=  0.5956863043975307
LogisticRegression(solver='newton-cg')

In [19]: from sklearn.ensemble import RandomForestClassifier
from sklearn import svm

In [20]: model = svm.SVC() #select the algorithm
model.fit(n_X_train,Y_train) # we train the algorithm with the training data and the training output
prediction=model.predict(n_X_test) #now we pass the testing data to the trained algorithm
print('The accuracy of the SVM is:',accuracy_score(prediction,Y_test))#now we check the accuracy of the algorithm.
#we pass the predicted output by the model and the actual output

The accuracy of the SVM is: 0.5849401722429692

In [21]: rforest=RandomForestClassifier()
rforest.fit(n_X_train,Y_train)
y_pred_rforest=rforest.predict(n_X_test)
print("accuracy Random Forest= ",accuracy_score(Y_test,y_pred_rforest))

accuracy Random Forest=  0.9939791174453166

In [ ]:
```